

Ethereum in 25 Minutes

A guided tour through the depths of the Merkle
underworld

Why ethereum?

- Seeming public consensus circa 2013: blockchains are useful for... stuff
- Not just money! Asset issuance, crowdfunding, domain registration, title registration, gambling, prediction markets, internet of things, voting, hundreds of applications!

Problem: most existing blockchain protocols were designed like this:



Or, at best, like this:



So... why not make a protocol that works like this?*



*Special thanks to Microsoft for sponsoring this conference

The concept



It's a blockchain... yawn

With... a few tiny additions

- A built-in programming language
- Two types of accounts
 - User accounts (controlled by private keys)
 - Contracts (controlled by code)
- Anyone can create an application with any rules by defining it as a contract

DNS: The Hello World of Ethereum

```
data domains[](owner, ip)

def register(addr):
    if not self.domains[addr].owner:
        self.domains[addr].owner = msg.sender

def set_ip(addr, ip):
    if self.domains[addr].owner == msg.sender:
        self.domains[addr].ip = ip
```


State vs History

- State = “current” information
 - Account balances
 - Nonces
 - Contract code and contract storage
- History = things that happened
 - Transactions
 - Receipts
- Currently, all “full” nodes store state, some store history and some do not store history

State

- State consists of key value mapping addresses to account objects
- Every account object contains 4 pieces of data:
 - Nonce
 - Balance
 - Code hash (code = empty string for private key-controlled accounts)
 - Storage trie root

Code execution

- Every transaction specifies a TO address it sends to (unless it's creating a contract)
- The TO address's code runs
- Code can:
 - Send ETH to other contracts
 - Read/write storage
 - Call (ie. start execution in) other contracts

Code execution

- Every (full) node on the blockchain processes every transaction and stores the entire state, just like Bitcoin
 - It's bold because it's important.

Gas

- Halting problem
 - Cannot tell whether or not a program will run infinitely
- Solution: charge fee per computational step (“gas”)
- Special gas fees also applied to operations that take up storage

Gas limit

- Counterpart to the block size limit in bitcoin
- Voting mechanism (can upvote/downvote gas limit by 0.0976% every block)
- Default strategy
 - 4712388 minimum
 - Target 150% of long-term EMA of gas usage

Transactions

- **nonce** (anti-replay-attack)
- **gasprice** (amount of ether per unit gas)
- **startgas** (maximum gas consumable)
- **to** (destination address)
- **value** (amount of ETH to send)
- **data** (readable by contract code)
- **v, r, s** (ECDSA signature values)

Receipts

- Intermediate state root
- Cumulative gas used
- Logs

Logs

- Append-only, not readable by contracts
- ~10x cheaper than storage
- up to 4 “topics” (32 bytes), plus data
- Intended to allow efficient light client access to event records (eg. domain name changed, transaction occurred, etc)
- Bloom filter protocol to allow easier searching by topic

Ethereum Virtual Machine

- Stack
- Memory
- Storage
- Environment variables
- Logs
- Sub-calling

High-level languages

- Compile to EVM code
- Serpent
- Solidity
- LLL

The ABI

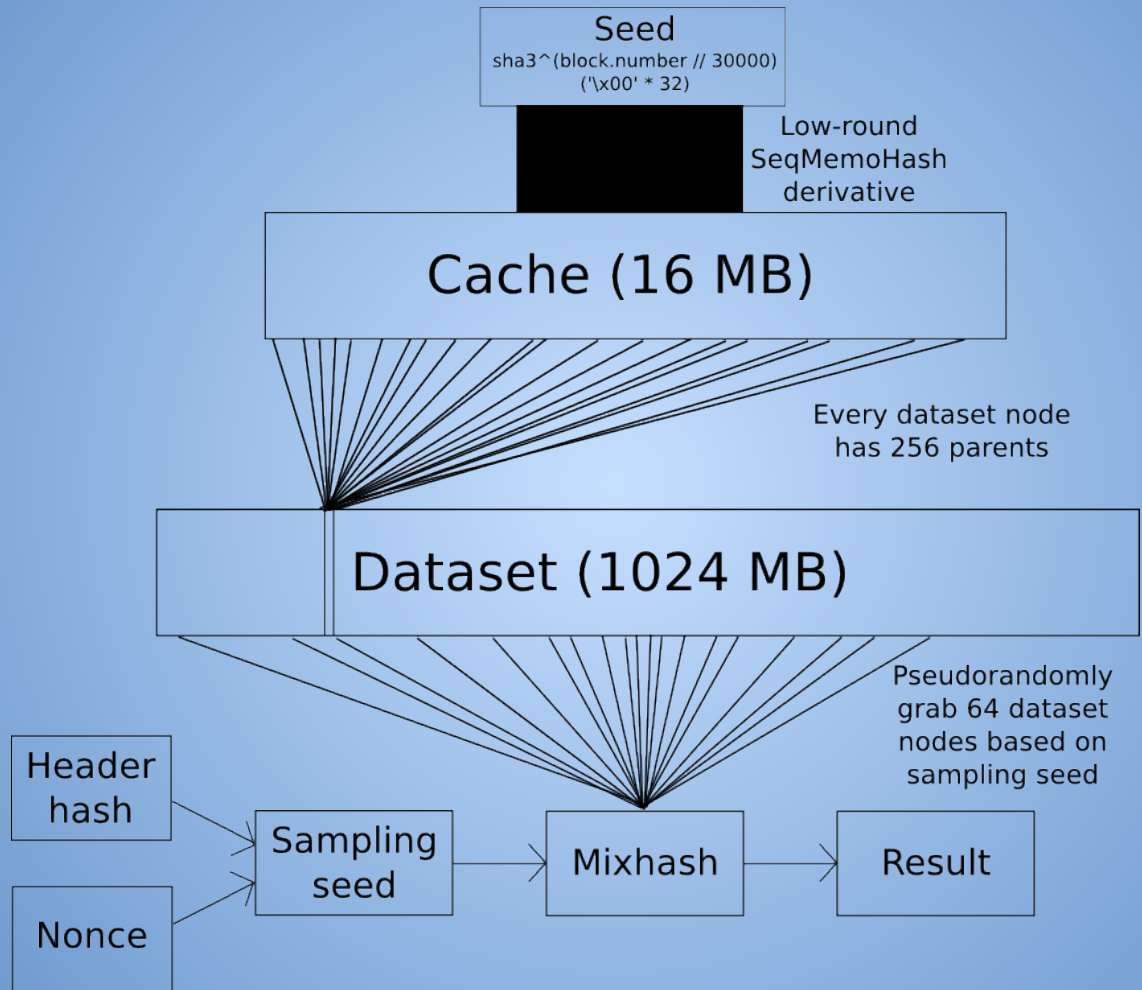
- Function calls are compiled into transaction data

[illegible]

- First 4 bytes: function ID
- Next 32 bytes: first argument ("foo.eth")
- Next 32 bytes: second argument (18 52 84 120 -> 0x12345678)

A cool new mining algorithm

- Goal: GPU-friendly, ASIC-hard
- CPU-friendly too risky a target, GPU “safe” balance, also limits botnet/AWS concerns
- Uses memory-hardness to achieve this goal
- Uses a multi-level DAG construction to achieve this with light-client friendliness




A cool new mining algorithm

- **Full node mining strategy:** re-build the full DAG every 30000 blocks, quickly compute mixhashes
- **Light node verification strategy:** re-build only the cache every 30000 blocks, more slowly compute mixhash once by computing 256 DAG nodes on the fly
 - Less efficient, but we only need to do this once


Fast block times

- The problem with fast block times: stale rates due to network latency




Normally, this leads to centralization risks

- Last block zero stale rate effect
- Network connectivity economies of scale



Our solution: uncles



- Our solution: allow stale blocks to be re-included and receive most of the reward
- Cuts centralization incentives by ~85%



Merkle trees

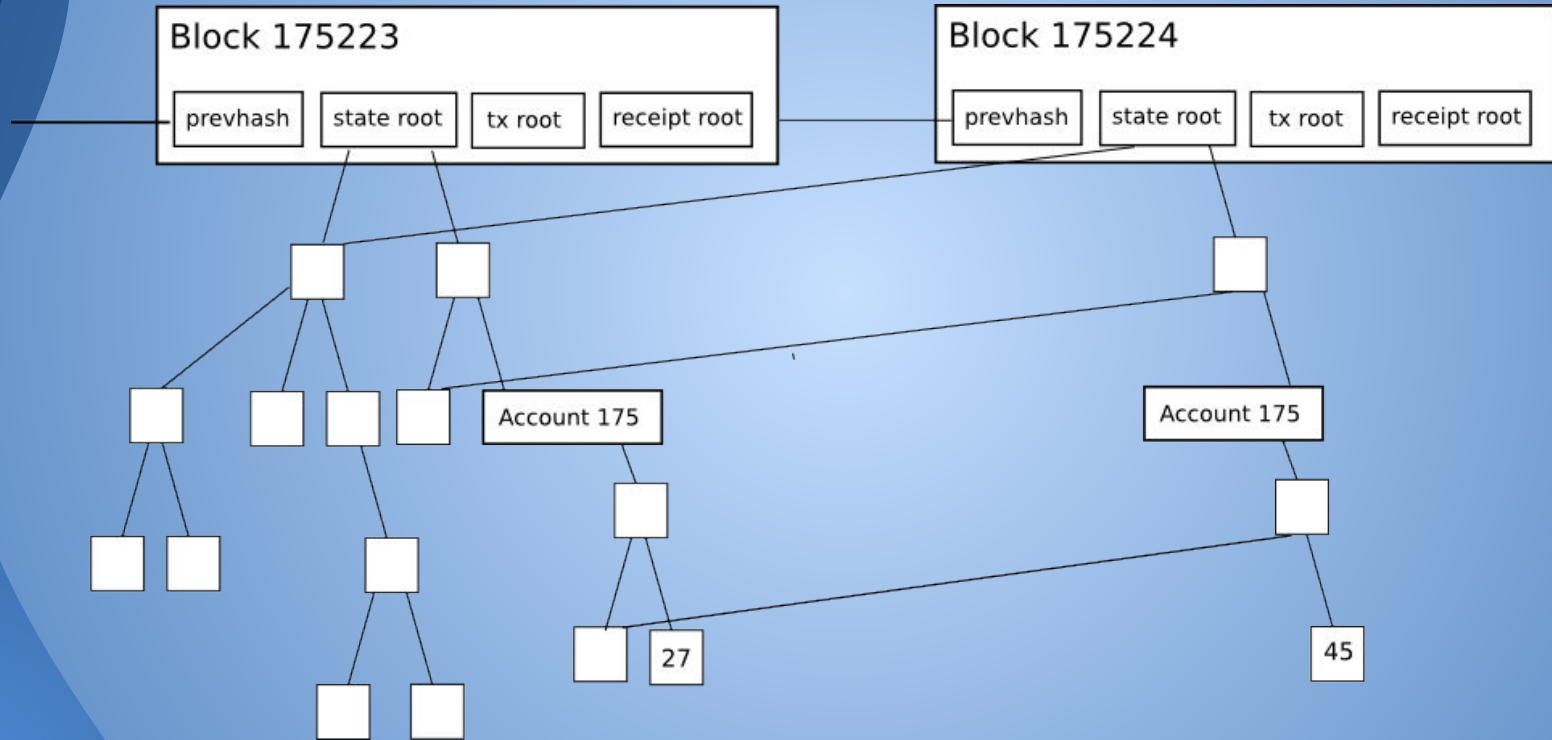


Merkle trees



Allow for efficiently verifiable proofs that a transaction was included in a block

Merkle trees in Ethereum



Merkle trees in Ethereum

- Every block header contains three tries
 - Transactions
 - State
 - Receipts
- Patricia trees used in order to allow efficient insert/delete operations

Future directions

- PoS
- Privacy support (ZK-snark precompiles, ring sigs, etc)
- Blockchain rent
- VM upgrades
- More flexible use of storage
- Scalability